

# Reinforcement learning lab report

Adrian-Gabriel Bălănescu

June 18, 2021

# Contents

<b>1</b>	<b>Game description</b>	<b>2</b>
1.1	Gameplay . . . . .	2
1.2	Gravitar in SLM-lab . . . . .	3
<b>2</b>	<b>Algorithms</b>	<b>4</b>
2.1	Reinforce . . . . .	4
2.2	DQN - Deep Q-Network . . . . .	4
2.3	A2C - Advantage Actor-Critic . . . . .	4
2.3.1	Estimating the advantage with n-Step returns . . . . .	5
2.3.2	Estimating the advantage with Generalized Advantage Estimation . . . . .	5
2.4	PPO - Proximal Policy Optimization . . . . .	5
<b>3</b>	<b>Experimental Results</b>	<b>6</b>
<b>4</b>	<b>Conclusions</b>	<b>7</b>
	<b>Bibliography</b>	<b>8</b>

# Chapter 1

## Game description

*Gravitar* is an arcade style game released by Atari in 1982. The game was known for its high level of difficulty [1].

### 1.1 Gameplay

The game takes place in a fictional solar system. The player controls a spacecraft with the help of five buttons: two to rotate the ship left and right, one to shoot, one to activate the thruster and one for both a tractor beam and a force field. In the side-view levels, the player has to destroy red bunkers that shoot constantly, and can also use the tractor beam to pick up blue fuel tanks. Once all of the bunkers are destroyed, the planet will blow up, and the player will earn a bonus. Once all planets are destroyed, the player will move onto another solar system. The player will lose a life if he crashes into the terrain or gets hit by an enemy's shot, and the game will end immediately if fuel runs out [1].

After completing all 11 planets (or alternatively completing the reactor three times) the player enters the second universe and the gravity will reverse. Instead of dragging the ship towards the planet surface, the gravity pushes it away. In the third universe the landscape becomes invisible and the gravity is positive again. The final, fourth universe, has invisible landscape and reverse gravity. After completing the fourth universe the game starts over [1].



Figure 1.1: Gravitar logo



Figure 1.2: A level from Gravitar




Figure 1.3: The reactor

## 1.2 Gravitar in SLM-lab

*SLM-lab* [2] is a software framework for reproducible reinforcement learning (RL) research. It implements most of the canonical RL algorithms in a modular platform that allows for easy experimentation and analysis. It includes the OpenAI Gym [3] environments, which contains a lot of Atari games emulations, one being Gravitar.

In this environment, the observation is an RGB image of the screen, which is an array of shape (210, 160, 3). Each action is repeatedly performed for a duration of  $k$  frames, where  $k$  is uniformly sampled from  $\{2, 3, 4\}$ . The action space is discrete and has 18 available commands like fire, up, down, right, left and so on.

The Gravitar environment can be loaded in SLM-lab by specifying in the .json configuration file, at the "env" section the "name" of the environment. In this case, "GravitarNoFrameskip-v4".

# Chapter 2

## Algorithms

In this chapter, the algorithms used will be briefly explained.

### 2.1 Reinforce

Reinforce is the most basic policy-gradient algorithm.

### 2.2 DQN - Deep Q-Network

...

### 2.3 A2C - Advantage Actor-Critic

The Advantage Actor-Critic algorithm relies on two important concepts:

**The Actor-Critic architecture.** This type of algorithms learn two parametrized functions simultaneously. The actor learns the actual policy that the agent uses to act in the world, while the critic learns the value function that evaluates the state-action pairs.

Using the critic to provide a reinforcement signal has the advantage of being able to feed a more continuous signal for training that is useful for example in sparse rewards environments. However, training becomes more difficult because the policy is dependent on the quality of the value function. In the beginning of the training process, the value function is not learned, so the evaluation signal provided doesn't help much to select better actions. Therefore, usually an Advantage function is actually learned.

**Reinforcing signal based on the Advantage function.** The advantage function measures how much an action is better or worse than a policy's average action in a certain state. It is defined as follows:

$$A^{\pi}(s_t, a_t) = Q^{\pi}(s_t, a_t) - V^{\pi}(s_t) \quad (2.1)$$

It has the following properties:

- The advantage is a relative measure. It quantifies whether a particular action  $a$  in a state  $s$  will improve the policy relative to the value  $V^\pi(s)$ .
- TODO

### **2.3.1 Estimating the advantage with n-Step returns**

### **2.3.2 Estimating the advantage with Generalized Advantage Estimation**

## **2.4 PPO - Proximal Policy Optimization**

PPO aims at avoiding the problem of performance collapse through a surrogate objective that is guaranteed to have monotonic improvement.

## Chapter 3

# Experimental Results

(here please provide, for each algorithm, the full config file, the description of the config, similar to those in the course, and the graphs of the results; try to train each algorithm for at least 500k-1M steps; make sure to train all the algorithms for the same number of steps)

## **Chapter 4**

### **Conclusions**

(here please discuss the results obtained and compare the performance of the algorithms)



# Bibliography

- [1] *Gravitar*. <https://en.wikipedia.org/w/index.php?title=Gravitar&oldid=1018480060>, April 2021. Page Version ID: 1018480060. 1, 1.1
- [2] Wah Loon Keng and Laura Graesser. Slm lab. <https://github.com/kengz/SLM-Lab>, 2017. 1.2
- [3] Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. Openai gym. *arXiv preprint arXiv:1606.01540*, 2016. 1.2